Applied Data Science 2

DR ASHLEY SPINDLER

SEMESTER B, 2022

Lecture 9 - Outcomes

- Understand the principles of GPU Acceleration
- Explore computational bottlenecks, and ways to mitigate them

GPUs

What are GPUs?

- GPU – Graphics Processing Unit
- Most often related to videos games, photography and video production
- Often called the “soul” of a PC
- Now form the backbone of ML development

CPU vs GPU

CPU – Central Processing Unit

- Composed of a handful of "cores"
- Low Latency
- Perform serial calculations
- Can only perform a few tasks at once

GPU Graphics Processing Unit

- Large number of cores
- High memory throughput
- Performs parallel calculations
- Can perform 1000s of tasks at once

CPU vs GPU

CPU – Intel i9 series

- ~800 GFLOPS calculation speed
- 18 cores

GPU – NVIDIA Tesla A100

- 10-300 TFLOPS calculation speed
- Equivalent of 6912 cores

GPUs are specialists

- GPUs are highly specialised in performing very specific, often very simple calculations incredibly fast
- E.g. multiplying together matrices, performing convolutions—the GPU can perform thousands of these a second

GPUs are ideal for ML

- A neural network is a graph of often very simple calculations—a Dense layer is just a matrix multiplication and addition—which is performed many, many times

- Even simple models can achieve massive speed ups compared to CPUs—e.g. the MLP from tutorial 8 runs 3x faster on a Tesla K80, compared to CPU

Multiple GPUs can be used together

- Because GPUs perform operations in parallel by design, you can link up multiple GPUs together and increase calculation speed
- Not always beneficial, as you can hit slow downs from data transfer between GPUs

GPU Summary

- GPUs utilise specialist hardware and algorithms to perform many thousands of simple calculations at the same time
- This makes them ideal for neural networks, which require many applications of matrix multiplications, additions, and other similar operations
- TensorFlow will automatically use a GPU if it is available, but we can also use multiple GPUs at once

Bottlenecks

In a perfect world…

- We would have very high utilisation of the CPU and GPU resources
- Our input pipelines would efficiently load, transform and move our data from the storage onto the CPU and GPU
- There would be no lag between GPU operations finishing on one batch of data, and the next batch arriving

Computation Pipelines

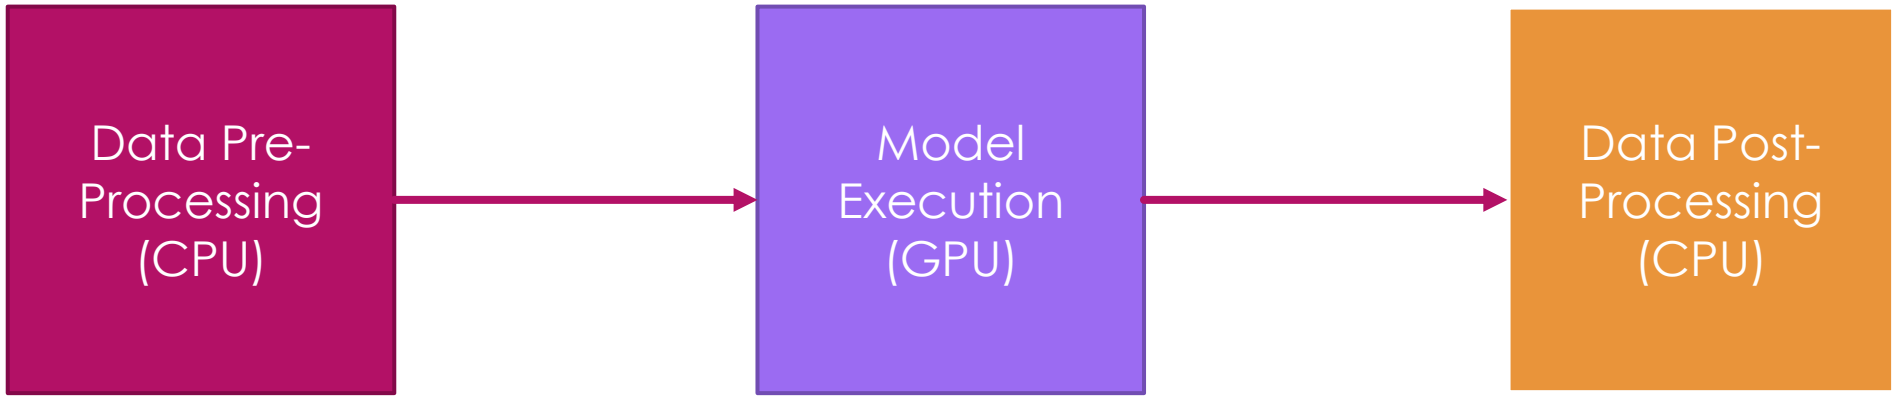

Computation Pipelines

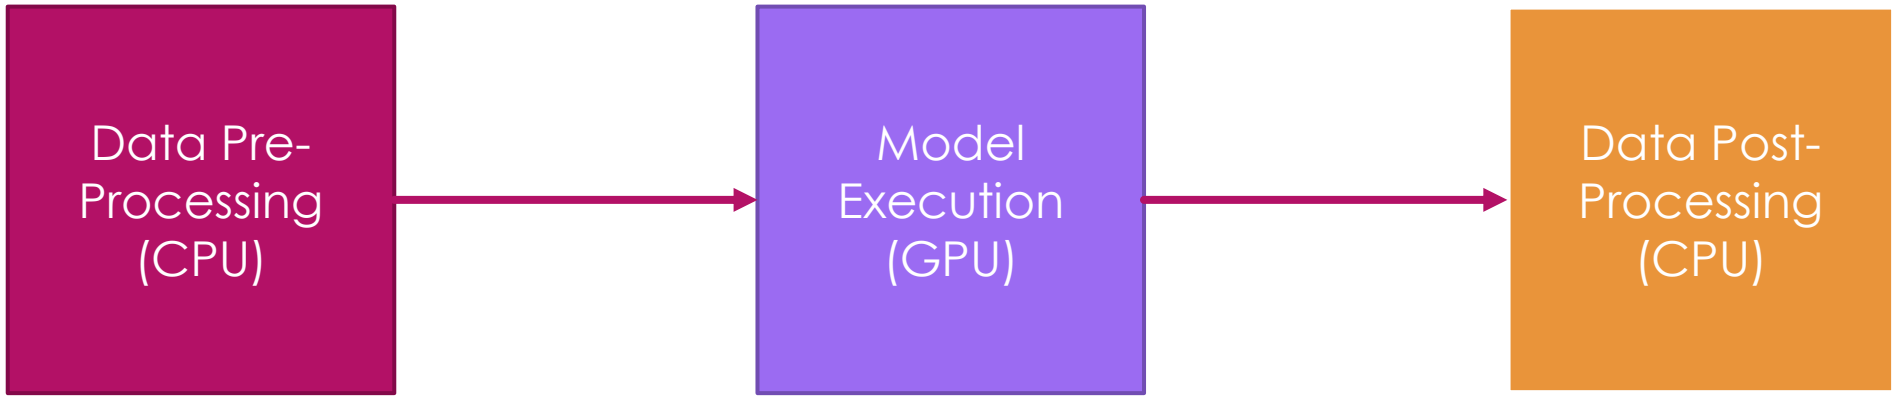

Reading data from HD
Processing inputs
Batching and Shuffling

Some pre-processing
Passing data through model layers
Calculating Losses, Metrics and Gradients
Updating Model Weights

Storing Losses and Metrics
Saving Checkpoints
Visualisation and Analysis

Computation Pipelines

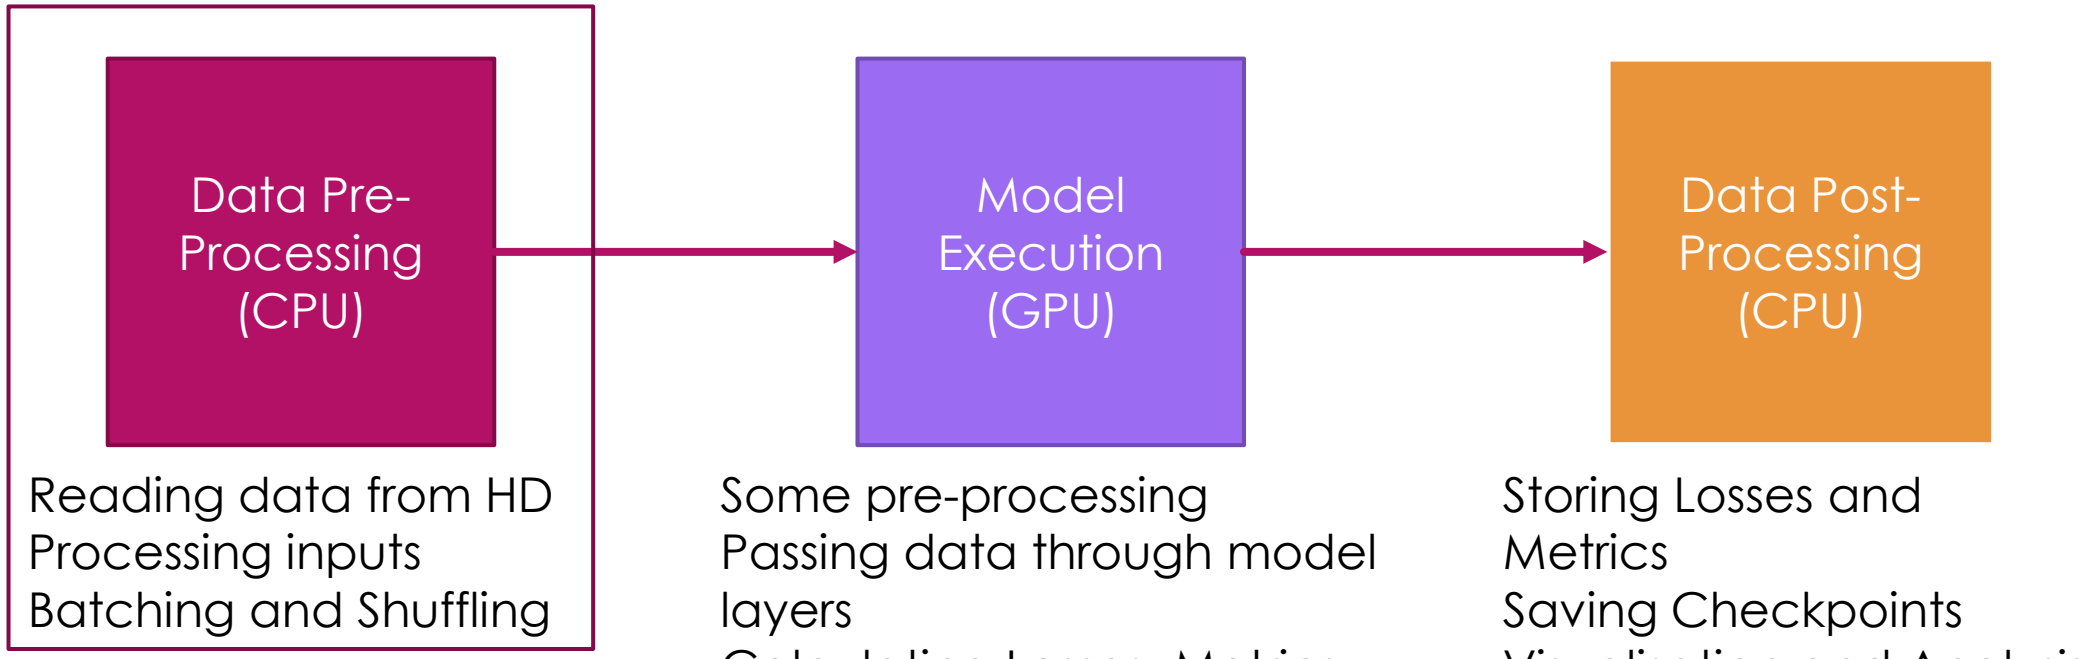

Reading data from HD
Processing inputs
Batching and Shuffling

Most Common Bottleneck is in CPU Pre-Processing

Some pre-processing
Passing data through model layers
Calculating Losses, Metrics and Gradients
Updating Model Weights

Storing Losses and Metrics
Saving Checkpoints
Visualisation and Analysis

Processing Bottlenecks

- Idle time on GPU while CPU is fetching data, likewise CPU idle when GPU is running
- Serial data extraction where only a single file is loaded at a time
- Data transformations performed sequentially on data samples
- Repeating calculations each training step
- Transforming an entire dataset instead of individual batches

Processing Bottlenecks

- Prefetch the next batch of data while model is executing
- Extract data in parallel instead of in sequentially
- Parallelise data transformations
- Cache the results of time consuming transformations
- Batch datasets before performing transformations

TensorBoard Profiler Demo

Further Reading

- TensorFlow Data Performance Guide: https://www.tensorflow.org/guide/data_performance
- TensorBoard Profiler: https://www.tensorflow.org/guide/profiler
- TensorFlow GPU Performance: https://www.tensorflow.org/guide/gpu_performance_analysis

Questions???